

INTERESTS	Distributed Systems, Key-Value Stores, Cloud Computing, Virtualization		
EDUCATION	<b>University of Utah</b> ..... Salt Lake City, USA Doctor of Philosophy in Computer Science, Ongoing, Advised by Prof. Ryan Stutsman		
PUBLICATIONS	Achieving High Throughput and Elasticity in a Larger-than-Memory Store <b>VLDB 2021</b> <b>Chinmay Kulkarni</b> , Badrish Chandramouli, and Ryan Stutsman     (Under Revision)		
	Adaptive Placement for In-memory Storage Functions <b>ATC 2020</b> Ankit Bhardwaj, <b>Chinmay Kulkarni</b> , and Ryan Stutsman		
	Splinter: Bare-Metal Extensions for Multi-Tenant Low-Latency Storage <b>OSDI 2018</b> <b>Chinmay Kulkarni</b> , Sara Moore, Mazhar Naqvi, Tian Zhang, Robert Ricci, and Ryan Stutsman		
	Rocksteady: Fast Migration for Low-latency In-memory Storage <b>SOSP 2017</b> <b>Chinmay Kulkarni</b> , Aniraj Kesavan, Tian Zhang, Robert Ricci, and Ryan Stutsman		
OPEN SOURCE	<b>microsoft/FASTER</b>	<b>vmware/node-replication</b>	<b>utah-scs/splinter</b>
EXPERIENCE	<b>University of Utah</b> ..... Salt Lake City, USA <i>Research Assistant advised by Ryan Stutsman, 2016 - Present</i> Worked on multi-tenant low-latency stores that can be extended at runtime using type- and memory-safe functions, and fast, low impact data migration protocols.		
	<b>Google</b> ..... Sunnyvale, USA <i>Research Intern hosted by Larry Kai, Summer 2020</i> Worked on defining and measuring the availability of Google services. Designed and built a dashboard that Google engineers can use to visualize and monitor the availability of their service.		
	<b>VMware</b> ..... Palo Alto, USA <i>Research Intern hosted by Gerd Zellweger, Summer 2019</i> Designed, built, tested and evaluated a Rust library that constructs a highly scalable, linearizable, concurrent data structure from a single threaded implementation.		
	<b>Microsoft</b> ..... Redmond, USA <i>Research Intern hosted by Badrish Chandramouli, Summer 2018</i> Worked on an RPC layer and scale out protocol for FASTER, a key-value store that scales linearly across cores to service 160 million updates per second.		
SERVICE	<b>HotCloud'20</b> , <i>External Reviewer</i>		
TALKS AND POSTERS	<b>Scaling an Operating System to Many Cores Using a System Call Log</b> <i>SOSP 2019 (Poster), Huntsville, Ontario, Canada</i>		
	<b>Raising The Efficiency of <math>\mu</math>Storage</b> <i>Google PhD Fellowship Summit 2019, Mountain View, California, USA</i>		
	<b>Splinter: Bare-Metal Extensions for Multi-Tenant Low-Latency Storage</b> <i>OSDI 2018, Carlsbad, California, USA</i>		
	<b>Rocksteady: Fast Migration for Low-latency In-memory Storage</b> <i>SOSP 2017, Shanghai, China</i>		
AWARDS	<b>Google PhD Fellowship</b> , <i>Systems and Networking, 2019</i>		
SKILLS	Rust, Python, R, C++, Kernel-bypass networking, Lock-free programming		