

# Achieving High Throughput and Elasticity in a Larger-than-Memory Store

**Chinmay Kulkarni**, Badrish Chandramouli\*, Ryan Stutsman  
University of Utah, \*Microsoft Research

# One Slide Summary

Multi-core optimized single-node key-value stores are emerging

→ **Very high throughput ~100 Million events/sec/server.** Ex: FASTER (SIGMOD'18)


**Problem:** Retain high throughput in the public cloud

- Avoid request dispatch and network bottlenecks
- Support reconfiguration/migration while preserving throughput


**Shadowfax:** Distributed key-value store built on FASTER

- 100 Million events/sec/VM **on Azure**
- **930 Million events/sec** on CloudLab cluster

# Seastar: State-of-the-art Cloud Key-Value Store



# Seastar: Records Partitioned Across Cores




**Hash partition  
records across vCPUs**

**Avoids  
synchronization on  
records**



**Cloud  
Network**

# Seastar: Records Partitioned Across Cores




**Dispatcher on every  
vCPU**

**Each listens on  
different TCP port**




**Cloud  
Network**

# Seastar: Requests Routed Within Server



Each dispatcher routes requests to correct partition

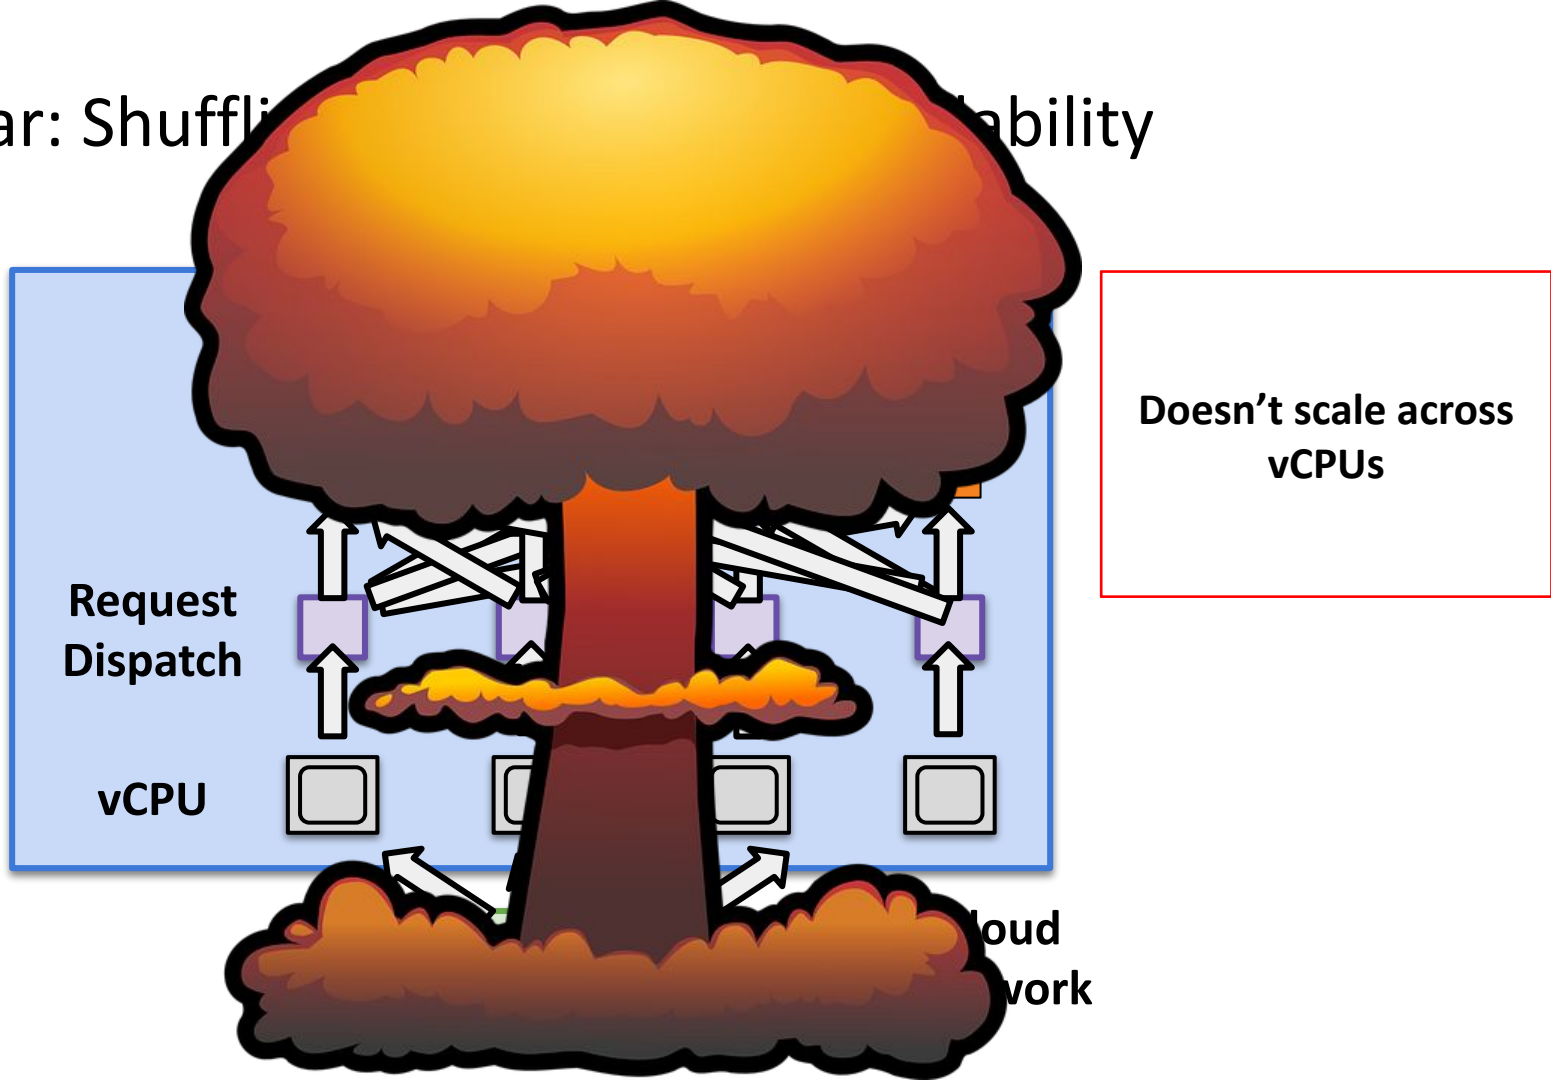
# Seastar: Shuffling Requests Limits Scalability



Each dispatcher  
routes requests to  
correct partition


Expensive shuffle!

# Seastar: Shuffling for Stability







# Seastar: Clients Could Route Requests Correctly



# Seastar: Clients Could Route Requests Correctly




# Seastar: Bad For Skewed Workloads




Bad for skewed  
workloads

# Seastar: Migration Has To Be Single Threaded



# Seastar: Head-of-line Blocking During Migration



# Shadowfax: Design

## **Partitioned Request Dispatch, Shared Data**

- vCPUs share lock-free index → record synchronization handled by cache coherence
- Migration can be parallelized, does not block requests on hash ranges

## **End-to-End Asynchronous Clients**

- Issue pipelined asynchronous batches of requests to amortize network overhead
- Good for workloads with inter-request independence. Ex: click counts, heartbeats etc.

# Shadowfax: Refer To Paper

## **Asynchronous Global Cuts**

- Per server view numbers represent hash ranges owned by server
- Cores avoid coordination during migration ownership changes

## **Scale-out and Hash Migration Protocol**

- Leverages asynchronous global cuts; is parallel and deadlock-free
- Fault-tolerant, can be cancelled; can recover if servers crash


## **Spans Main-memory, Local SSD, and Shared Remote Storage**

- During migration, only transfer (hot) data in main-memory
- Transfer pointers to rest (cold), lazily clean these up during compaction

# Shadowfax: Partitioned Request Dispatch Shared Data

**Problem:** Avoid cross-core coordination at server

**Solution:** Offload all coordination to hardware cache-coherence protocol



Shared memory is  
good for skew

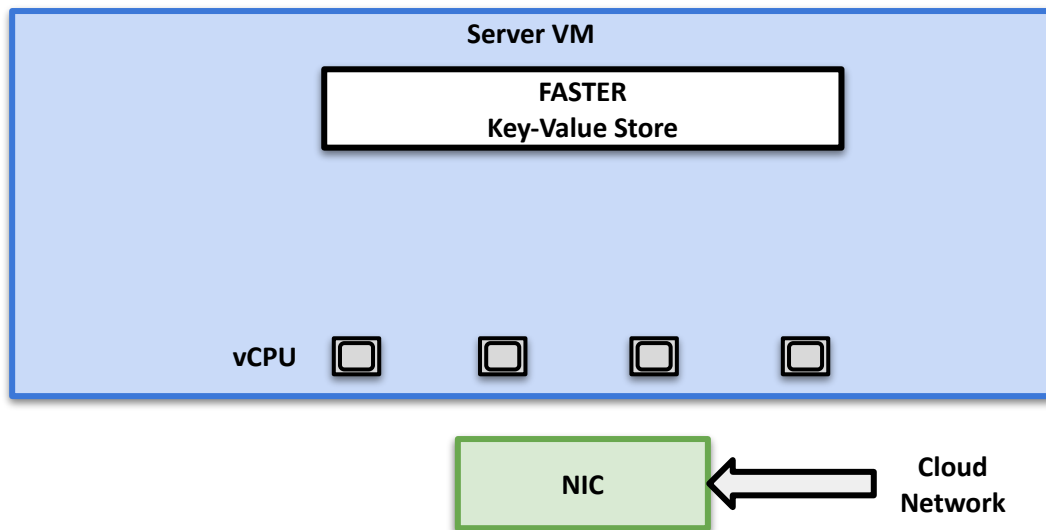
Can serve hot records  
in parallel



# Shadowfax: Partitioned Request Dispatch Shared Data

**Problem:** Avoid cross-core coordination at server

**Solution:** Offload all coordination to hardware cache-coherence protocol




**Lock-free index**

**Record synch happens  
at cache-coherence  
protocol (hw)**

# Shadowfax: Partitioned Request Dispatch Shared Data

**Problem:** Multi-core request processing requires cross-core coordination

**Solution:** Offload all coordination to hardware cache-coherence protocol




**Partition request  
dispatch**

**Each dispatcher listens  
on separate TCP port**

# Shadowfax: Partitioned Request Dispatch Shared Data

**Problem:** Keep servers saturated at index, not network or dispatch


**Solution:** Asynchronous client library, transparent network acceleration



# Shadowfax: Partitioned Request Dispatch Shared Data

**Problem:** Keep servers saturated at index, not network or dispatch

**Solution:** Asynchronous client library, transparent network acceleration




Client cores  
open  
*"sessions"*  
to  
server cores

# Shadowfax: Partitioned Request Dispatch Shared Data

**Problem:** Keep servers saturated at index, not network or dispatch

**Solution:** Asynchronous client library, transparent network acceleration




Requests are  
async

Enqueued in  
session  
buffers

# Shadowfax: Partitioned Request Dispatch Shared Data

**Problem:** Keep servers saturated at index, not network or dispatch

**Solution:** Asynchronous client library, transparent network acceleration




Pipelined  
batches are  
transmitted  
asynchronously

# Shadowfax: Partitioned Request Dispatch Shared Data

**Problem:** Keep servers saturated at index, not network or dispatch

**Solution:** Asynchronous client library, transparent network acceleration





Pipelined  
batches are  
transmitted  
asynchronously

# Shadowfax: Transparent Cloud Acceleration

**Problem:** Keep servers saturated at index, not network or dispatch

**Solution:** Asynchronous client library, transparent network acceleration




Hardware  
accelerated  
transport  
(TCP & Infr)



# Performance of Shadowfax


**YCSB-F (Read-Modify-Writes, Benchmark Ingest of Events), 250 Million objects**



**Saturates server at Index**

# Performance of Shadowfax

**YCSB-F (Read-Modify-Writes, Benchmark Ingest of Events), 250 Million objects**



**Saturates server at Index, 8.5x state-of-the-art**

# One Slide Summary

Multi-core optimized single-node key-value stores are emerging

→ **Very** high throughput ~**100 Million events/sec/server**. Ex: FASTER (SIGMOD'18)

**Problem:** Retain high throughput in the public cloud

- Avoid request dispatch and network bottlenecks
- Support reconfiguration/migration while preserving throughput

**Shadowfax:** Distributed key-value store built on FASTER

- 100 Million events/sec/VM **on Azure**
- **930 Million events/sec** on CloudLab cluster